



RV College of Engineering®

Mysore Road, RV Vidyaniketan Post, Bengaluru - 560059, Karnataka, India

AUTOMATED MULTI-AGENT GUI ORCHESTRATOR

AIML-PROJECT REPORT

Submitted by

Aditya Ankanath TR 1RV23IS007

Ankit Pathak 1RV23IS017

Under the guidance of

Dr. Merin Meleet

Associate Professor

Dept. of ISE

RV College of Engineering

In partial fulfillment for the award of degree of

Bachelor of Engineering

in

Information Science and Engineering

2025-2026

RV COLLEGE OF ENGINEERING[®], BENGALURU - 560059
(Autonomous Institution Affiliated to VTU, Belagavi)

DEPARTMENT OF INFORMATION SCIENCE AND ENGINEERING



CERTIFICATE

Certified that the project work titled “**Automated Multi-Agent GUI Orchestrator**” is carried out by **Aditya Ankanath TR (1RV23IS007)** and **Ankit Pathak (1RV23IS017)**, who are Bonafide student of R.V College of Engineering, Bangalore, in partial fulfillment for the award of degree of **Bachelor of Engineering in Information Science and engineering** of the Visvesvaraya Technological University, Belgaum during the year **2025-26**. It is certified that all corrections/suggestions indicated for the internal Assessment have been incorporated in the report deposited in the departmental library. The project report has been approved as it satisfies the academic requirements in respect of project work as a part of the course “Artificial Intelligence and Machine Learning” prescribed by the institution for the said degree.

Faculty in Charge

Head of the Department

External Evaluation

Name of Examiners

Signature with date

1

2



RV College of Engineering[®]

Mysore Road, RV Vidyaniketan Post, Bengaluru - 560059, Karnataka, India

DECLARATION

We, **Aditya Ankanath TR** and **Ankit Pathak** students of fifth semester B.E., Department of Information Science and Engineering, RV College of Engineering, Bengaluru, hereby declare that the project titled "**Automated Multi-Agent GUI Orchestrator**" has been carried out by us and submitted in partial fulfilment for the award of degree of **Bachelor of Engineering** in **Information Science and Engineering** during the year 2025-2026

Further we declare that the content of the dissertation has not been submitted previously by anybody for the award of any degree or diploma to any other university. We also declare that any Intellectual property rights generated out of this project carried out at RVCE will be the property of RV College of Engineering, Bengaluru and we will be among the authors of the same.

Place: Bengaluru

Date:

Signature

ACKNOWLEDGEMENTS

We would like to express our sincere gratitude to our project guide, Dr. Merin Meleet, for their invaluable guidance, constant encouragement, and dedicated support throughout the course of this project. Their insights and mentorship played a crucial role in shaping our work and helping us overcome challenges at every stage.

We extend our heartfelt thanks to Dr. Mamatha G S, Head of the Department of Information Technology, and to all the respected faculty members for their continuous assistance, motivation, and constructive feedback. We are also grateful to the teaching and non-teaching staff of the Computer Department for their timely support in all aspects. Our sincere appreciation goes to the Librarian for providing us with the necessary reference materials and resources essential for this project.

Finally, we express our deepest gratitude to our parents for their unwavering support and encouragement, and to our friends for their constant motivation, positivity, and companionship throughout this academic journey. Their support made this experience not only productive but also truly memorable.

ABSTRACT

Modern human–computer interaction systems increasingly aim for autonomy and efficiency, yet they often struggle with reliable visual understanding of real-world user interfaces. Conventional automation tools depend heavily on rigid scripts, fixed coordinates, or accessibility trees, making them brittle to UI changes, screen resolution variations, and dynamic layouts. While large vision-language models can interpret screens, they are typically over-parameterized, slow, and unsuitable for real-time desktop interaction. Additionally, most existing approaches mix reasoning, planning, and execution within a single heavy model, reducing controllability and increasing latency, which limits their practical deployment for lightweight computer-use automation.

To address these challenges, the proposed system introduces a lightweight AI-driven visual grounding framework for desktop automation that separates reasoning, grounding, and execution into specialized components. The methodology leverages a locally deployed LLM to reinterpret user commands into precise visual descriptions, followed by a fine-tuned grounding agent that operates solely on screenshots to predict exact bounding box coordinates of target UI elements. Instead of performing complex multi-task reasoning, the grounding model is optimized to output only pixel-level coordinates, significantly reducing computational overhead. The system supports diverse UI contexts, including taskbars, browser windows, and application interfaces, and adapts dynamically to varying screen resolutions. A small, curated dataset of annotated screenshots is used for fine-tuning, avoiding large-scale data requirements while maintaining accuracy.

The final outcome of this project is a modular, efficient computer-use automation framework capable of translating natural language commands into precise on-screen actions through coordinate-based grounding. By decoupling command understanding from visual localization and restricting the grounding agent’s output to bounding boxes alone, the system achieves faster inference, improved robustness, and better reproducibility compared to monolithic vision-language agents. This approach enables reliable desktop interaction with minimal computational resources and provides a scalable foundation for future research in human-in-the-loop automation, intelligent agents, and assistive technologies. Ultimately, the solution demonstrates how lightweight fine-tuned grounding agents, guided by LLM-based query refinement, can enable practical and dependable UI automation in real-world environments.

CONTENTS

Abstract	i
List of Figures	v
List of Tables	vi
List of Abbreviations	vi
Publication Details	viii
1 Introduction	1
1.1 Terminology.....	1
1.2 Scope and relevance.....	2
1.3 Motivation.....	3
1.4 Problem Statement.....	4
1.5 Objectives.....	5
1.6 Summary.....	5
2 Literature Survey	6
2.1 Literature Review.....	8
2.2 Functional Requirements.....	9
2.3 Hardware Requirements.....	10
2.4 Software Requirements.....	10
2.5 Summary.....	10
3 Design of the System	11
3.1 Theory and Concepts	11
3.2 Dataset Description...	13
3.3 Design and Methodology.....	14
3.4 System Architecture.....	16
3.5 Tools and APIs...	16
3.6 Summary.....	16

4 Implementation and Testing	17
4.1 Implementation Requirements.....	18
4.2 Implementation Tool Features.....	19
4.3 Code Snippets with explanation.....	19
4.4 Testing.....	20
4.5 Summary.....	20
5 Results and Analysis	21
5.1 Results.....	21
5.2 Benchmarking and Analysis.....	21
5.3 Screenshots.....	21
5.4 Innovative Component.....	22
5.5 Summary.....	22
6 Conclusion	23
6.1 Conclusion.....	23
6.2 Limitations.....	23
6.3 Future Scope.....	23
References	24
Appendix	
A. Dataset sample	
B Research Paper (IEEE format)	

LIST OF FIGURES

Figure 3.1 System Architecture Design.....	21
Figure 3.2 Fine Tune Model with Custom Screenshots.....	25
Figure 3.2 Grounding Agent Output with Area of Interest region.....	26
Figure 3.2 Llama3.2 based NLP query expansion.....	26
Figure 5.1 Example Command-to-Open CMD using Automation Interface...	33
Appendix Figure 1 Hugging Face Dataset Card.....	43
Appendix Figure 2 Dataset Snapshots.....	44
Appendix Figure 3 Hugging Face Dataset Name and Reference.....	44

LIST OF TABLES

Figure 2.1 Comparison of Existing Vision Grounding and UI Localization Model	17
Figure 3.1 System Architecture Components and Functional Roles.....	24
Figure 5.1 Performance Evaluation of the Proposed Framework.....	31
Figure 5.2 Example Command-to-Bounding-Box Outputs.....	32

LIST OF ABBREVIATIONS

LLM – Large Language Model

VLM – Vision–Language Model

UI – User Interface

GPU – Graphics Processing Unit

CPU – Central Processing Unit

JSON – JavaScript Object Notation

HTTP – Hypertext Transfer Protocol

FPS – Frames Per Second

OCR – Optical Character Recognition

IoU – Intersection over Union

BBox – Bounding Box

CLI – Command Line Interface

RAG – Retrieval-Augmented Generation

CUDA – Compute Unified Device Architecture

FP16 – Half-Precision Floating Point

INT8 – 8-bit Integer Quantization

LLaMA – Large Language Model Meta AI

Qwen – Vision–Language Model family used for grounding

GroundingDINO – Open-vocabulary object grounding model

PyAutoGUI – Python-based GUI automation library

n8n – Workflow orchestration and automation platform

CHAPTER 1

INTRODUCTION

INTRODUCTION

The rapid evolution of artificial intelligence has significantly transformed the way users interact with computer systems. Traditional human–computer interaction relies on direct physical inputs such as mouse clicks and keyboard strokes, while modern automation systems aim to interpret high-level user commands and execute them autonomously. Recent advances in large language models and vision–language models have enabled machines to reason about visual content, understand user intent, and interact with graphical user interfaces. However, most existing systems prioritize general reasoning and textual explanations rather than precise, low-level control required for direct UI manipulation.

This project, titled “**Lightweight LLM-Assisted Visual Grounding for Desktop UI Automation**,” addresses this limitation by focusing exclusively on accurate coordinate prediction for UI elements. Instead of generating verbose reasoning or multi-modal outputs, the proposed system enforces a strict output format consisting of bounding box coordinates only. By separating language understanding from visual grounding, the framework achieves faster inference and greater reliability in real-world desktop environments. This approach bridges the gap between high-level user intent and low-level screen interaction, enabling efficient automation without excessive computational cost.

1.1 Terminology

- **Visual Grounding:** The process of mapping a textual description to a specific region within an image.
- **Bounding Box (BBox):** A rectangular region defined by four pixel coordinates representing the location of a UI element.
- **Vision–Language Model (VLM):** A neural model capable of jointly processing visual and textual inputs.
- **Large Language Model (LLM):** A transformer-based model used for understanding and expanding natural language commands.
- **Command Expansion:** The transformation of a short user command into a detailed visual description to improve grounding accuracy.
- **Coordinate-Only Output:** A strict output constraint where the model returns only numerical bounding box values.

- **UI Automation:** Programmatic control of graphical interfaces using predicted screen coordinates.
- **Screen Saliency:** The degree to which a region visually stands out and attracts model attention.
- **Normalized Coordinates:** Bounding box values scaled between 0 and 1 relative to screen dimensions.
- **Inference Latency:** The time taken by a model to produce an output from a given input.
- **Lightweight Model:** A model optimized for lower memory usage and faster execution.
- **Fine-Tuning:** The process of adapting a pretrained model to a specific task using a smaller dataset.

1.2 Scope and Relevance

The scope of this project is limited to single-screen desktop environments and static screenshots captured during normal system usage. The system focuses on identifying common UI elements such as application icons, search bars, buttons, and input fields using natural language commands. It does not aim to replace full-fledged autonomous agents or operating system-level automation tools, but instead proposes a lightweight grounding layer that can be integrated into existing automation pipelines.

The relevance of this work is significant in the context of real-time UI automation, accessibility tools, testing frameworks, and intelligent assistants. Heavy vision–language models often fail to meet latency and resource constraints required for continuous interaction. By emphasizing strict coordinate outputs and modular design, this project supports scalable deployment on consumer-grade hardware. It is particularly relevant for research in computer-use agents, assistive technologies, and intelligent automation systems.

1.3 Motivation

The primary motivation for this project arises from the practical limitations observed in existing vision–language automation systems. While state-of-the-art models demonstrate impressive reasoning capabilities, they frequently produce verbose explanations, inconsistent output formats, or ambiguous grounding results. Such behavior is unsuitable for downstream automation tasks that require precise and deterministic coordinates.

Additionally, deploying large end-to-end VLMs often demands high memory and compute resources, making them impractical for local or real-time applications. Through personal experimentation and system-level analysis, it became evident that separating language reasoning from visual localization could significantly reduce complexity while improving reliability. This project is motivated by the goal of creating a minimal, efficient, and research-friendly grounding system that prioritizes actionable outputs over unnecessary contextual

reasoning.

1.4 Problem Statement

Current vision–language models designed for computer interaction generate excessive contextual information and lack strict output constraints, making them inefficient for direct UI automation. Their high computational cost and unpredictable response formats hinder real-time deployment and integration with automation tools. Users require a system that can accurately identify UI elements and return precise coordinates without additional explanation or overhead. Furthermore, natural language commands are often ambiguous and insufficient for direct visual grounding. Without an intermediate semantic expansion step, grounding accuracy suffers significantly. This project addresses these challenges by introducing a two-stage framework that refines user intent through an LLM and enforces strict coordinate-only outputs from a lightweight vision grounding model.

1.5 Objectives

- **Design a coordinate-only grounding system:** The system is designed to output only bounding box coordinates for UI elements, eliminating unnecessary textual responses and ensuring deterministic automation.
- **Integrate LLM-based command expansion:** A local LLM is used to convert short commands into visually descriptive queries that improve grounding reliability.
- **Develop a lightweight vision grounding pipeline:** The vision model is optimized for fast inference and low memory usage, making it suitable for real-time desktop applications.
- **Enable fine-tuning with small datasets:** The framework supports fine-tuning using a limited number of annotated screenshots, avoiding the need for large-scale datasets.
- **Improve grounding accuracy and latency:** By decoupling reasoning and localization, the system aims to achieve better accuracy with reduced inference time.

1.6 Summary

This chapter introduced the motivation, scope, and significance of a lightweight LLM-assisted visual grounding system for desktop UI automation. It highlighted the limitations of existing vision–language models and established the need for strict coordinate-only outputs. The chapter concluded by outlining clear objectives that guide the design and development of an efficient and deployable grounding framework.

CHAPTER 2

LITERATURE SURVEY

2.1 Literature Review

ScreenAgent: Vision-Language Models for GUI Automation [1] presents an early and influential approach to desktop automation using multimodal reasoning. The authors propose a vision-language agent that interprets screenshots and natural language instructions to locate UI elements and perform actions. While the system demonstrates strong generalization across GUI tasks, its reliance on large multimodal models leads to high computational cost and slow inference. Additionally, bounding box localization lacks the precision required for real-time automation. Nevertheless, this work establishes a foundational framework for vision-based UI grounding and motivates the need for lightweight, coordinate-focused agents.

UIAct: Automated User Interface Interaction via Multimodal Learning [2] introduces a multimodal framework for automated UI navigation using joint visual–textual embeddings. The system achieves high task completion rates in controlled environments by predicting actions based on learned representations. However, it struggles with unseen UI layouts and does not enforce strict output constraints, often producing verbose or ambiguous results instead of precise coordinates. This limitation highlights the necessity for grounding models that output deterministic spatial information, directly influencing the design goals of our system.

Grounding DINO: Marrying DINO with Grounded Language

Understanding [3] proposes an open-vocabulary object grounding model capable of mapping textual queries to bounding boxes with high accuracy. Although highly effective for general object detection, its performance degrades when applied to fine-grained UI elements such as icons, buttons, and text fields.

Despite this limitation, Grounding DINO serves as a strong baseline for text-to-bounding-box grounding and informs the visual grounding layer of our architecture.

Visual Prompting for Desktop Automation [4] explores prompt-driven visual reasoning for interacting with desktop environments using screenshots and high-level commands. The study demonstrates that carefully designed visual prompts can guide models to perform complex UI interactions. However, the approach requires extensive prompt engineering and produces verbose outputs unsuitable for low-latency execution. This work emphasizes the trade-off between reasoning depth and execution efficiency, reinforcing the need for concise, coordinate-only outputs in automation systems.

Lightweight Vision-Language Models for Edge Deployment [5] focuses on optimizing multimodal models for deployment on resource-constrained hardware. The authors present techniques such as parameter reduction and task-specific fine-tuning to reduce inference latency. While these optimizations impact general reasoning capabilities, the study validates the effectiveness of lightweight, task-specialized models. This directly supports our decision to fine-tune a compact grounding model dedicated exclusively to coordinate prediction.

Command Understanding and Rewriting using Large Language Models [6] demonstrates that LLMs can effectively rewrite ambiguous or incomplete user commands into explicit task descriptions. Although the work does not integrate visual feedback, it establishes the effectiveness of LLMs for command refinement. This aligns directly with our use of a local LLaMA model to expand high-level automation commands into visually grounded descriptions suitable for downstream vision models.

End-to-End Multimodal Agents for Computer Use [7] presents a comprehensive framework that combines reasoning LLMs with vision models to enable autonomous computer interaction. While the system achieves impressive

task autonomy, the models are extremely large and unsuitable for local deployment or real-time execution. This limitation motivates our modular design, where reasoning and grounding are separated into efficient, independently optimized components.

UI Element Detection using Vision Transformers [8] investigates transformer-based vision backbones for detecting UI components. The approach achieves strong localization accuracy but requires large annotated datasets and does not support natural language queries. Despite these constraints, the work confirms the feasibility of fine-tuning vision models specifically for UI element localization, supporting our dataset-driven fine-tuning strategy.

Multimodal Command Execution in Desktop Environments [9] combines speech, text, and vision to perform desktop actions using bounding-box-based interaction. The authors note that errors in command interpretation often propagate into execution failures. This finding reinforces the importance of reliable command expansion prior to visual grounding, validating our two-stage pipeline design.

Precision-Grounded Visual Agents for Human-Computer Interaction [10] proposes grounding agents that output only spatial coordinates for downstream automation. By eliminating extraneous reasoning text, the approach achieves improved execution reliability and lower latency. This work closely aligns with our core objective of producing a strict bounding-box-only output, serving as a direct conceptual reference for our system.

Model Name	Model Type	Parameter Size	Output Capability	Limitations
Qwen2.5-VL-72B	Vision-Language Model	~72B	Text + Box + Reasoning	Heavy, slow
Qwen2.5-VL-7B	Vision-Language Model	~7B	Text + Box	Verbose output
Grounding DINO	Vision Detector	~300M	Multiple boxes	Needs thresholds
ScaleCUA-3B	UI-focused VLM	~4B	Box + Action	Planner coupled
YOLOv8-UI	Object Detector	~100M	Class boxes	No text grounding
Proposed System	LLM + Lightweight VLM	~7B + local LLM	Single bbox only	Limited semantics

Tab 2.1 Comparison of Existing Vision Grounding and UI Localization Models

2.2 Functional Requirements

- FR1: Screenshot Input Handling**

Allow the system to accept desktop screenshots as primary visual input.

- FR2: Natural Language Command Input**

Enable users to issue high-level natural language commands describing desired UI actions.

- FR3: Command Expansion using LLM**

Use a lightweight LLM to rewrite ambiguous commands into explicit visual descriptions.

- FR4: Visual Grounding and Localization**

Identify the target UI element in the screenshot and compute precise bounding box coordinates.

• FR5: Coordinate-Only Output Generation

Ensure the system outputs only bounding box coordinates without additional text.

• FR6: Full-Screen Grounding Support

Support grounding across the entire desktop, not limited to specific UI regions.

• FR7: Model Inference Optimization

Maintain low-latency inference suitable for real-time automation.

• FR8: Failure Detection and Null Output

Return a null bounding box when the target element is not visible.

• FR9: Dataset Management for Fine-Tuning

Store and manage screenshot-command-coordinate triplets for supervised fine-tuning.

• FR10: Fine-Tuned Model Deployment

Support loading and execution of fine-tuned grounding models locally.

• FR11: Deterministic Output Enforcement

Guarantee structured and deterministic output formats for automation safety.

• FR12: Hardware-Aware Execution

Adapt model execution based on available CPU/GPU resources.

• FR13: Logging and Analysis Support

Log expanded queries, predicted boxes, and confidence scores for evaluation.

• FR14: Scalability of Command Vocabulary

Support gradual expansion of supported UI commands without retraining the entire system.

• FR15: Robust Error Handling

Gracefully handle invalid inputs, corrupted images, or inference failures.

2.3 Hardware Requirements

Inference and Processing System:

- Processor: Quad-core CPU (GPU optional)
- Memory: 8GB RAM (16GB recommended)
- Storage: 50GB SSD
- Network: Optional (offline capable)
- OS: Windows / Linux

Development and Training System:

- Processor: Multi-core CPU with CUDA-enabled GPU
 - Memory: 16GB RAM (32GB preferred)
 - Storage: 200GB SSD
 - OS: Linux (Ubuntu preferred)
-

2.4 Software Requirements

Core Frameworks & Libraries:

- PyTorch
- Hugging Face Transformers
- OpenCV / PIL
- NumPy
- Qwen2.5-VL
- LLaMA 3.2 (Local)

Programming Languages:

- Python

Development Tools:

- Jupyter Notebook
 - VS Code
 - Git & GitHub
-

2.5 Summary

Existing research demonstrates the feasibility of vision-language agents for desktop automation and UI grounding. However, most systems rely on large, slow models that produce verbose outputs unsuitable for precise automation. The reviewed studies highlight gaps in efficiency, determinism, and coordinate accuracy. These limitations motivate the proposed lightweight, fine-tuned grounding framework

CHAPTER 3

DESIGN OF THE SYSTEM

3.1 Theory and Concepts

This section describes the core technologies and theoretical foundations underlying the lightweight visual grounding and desktop automation framework.

Visual Grounding Theory

Visual grounding forms the core foundation of UI element localization in this project. A desktop screen is treated as a structured visual space where interactive UI components such as icons, buttons, input fields, and taskbar elements are considered target entities. By mapping natural language queries to precise pixel-level bounding boxes, visual grounding enables accurate interaction with graphical user interfaces. This approach allows deterministic execution of user commands by directly linking semantic intent to screen coordinates.

Vision–Language Models (VLMs)

Vision–Language Models are used to associate textual descriptions with visual patterns present in screenshots. Unlike traditional object detectors that rely on fixed labels, VLMs understand descriptive phrases such as color, shape, relative position, and contextual cues. In this project, the VLM is responsible solely for predicting a single bounding box corresponding to a queried UI element, enabling minimal and focused inference without unnecessary reasoning overhead.

Language Model–Driven Query Expansion

Large Language Models are employed to transform high-level user commands into visually descriptive grounding queries. Commands such as “open chrome” or “search youtube” are expanded into concrete visual attributes like shape, color, iconography, and expected screen region. This intermediate reasoning step simplifies the task for the VLM and significantly improves grounding accuracy without requiring large-scale retraining.

Lightweight Inference Design

The system is optimized for lightweight inference by restricting model outputs strictly to bounding box coordinates. All explanatory text, reasoning chains, and auxiliary metadata are eliminated from the inference pipeline. This design reduces latency, memory usage, and computational cost while maintaining precise UI interaction capabilities suitable for real-time

desktop automation.

Bounding Box Representation

Bounding boxes are represented using absolute pixel coordinates in the format (x_1, y_1, x_2, y_2) , corresponding to the top-left and bottom-right corners of the detected UI element. This representation allows seamless integration with automation tools such as mouse control and keyboard simulation libraries. Normalized coordinates may optionally be derived for logging and evaluation purposes.

Saliency-Based Visual Localization

The system implicitly leverages visual saliency by focusing on high-contrast, frequently interacted UI regions such as taskbars, address bars, and content panels. Saliency-aware grounding improves robustness when multiple visually similar elements exist on screen, ensuring that the most interaction-relevant component is selected.

Desktop Automation Integration

Once bounding box coordinates are obtained, the system executes deterministic actions such as click, type, or focus using OS-level automation interfaces. This decoupling of perception (VLM) and action (automation engine) ensures modularity and reliable execution across different desktop states.

Error-Constrained Output Enforcement

Strict output validation is enforced to guarantee that the model returns either a valid bounding box or a null response. Any deviation from the expected JSON structure is rejected, ensuring predictable downstream execution and preventing unintended UI interactions.

3.2 Dataset Description

Desktop UI Screenshot Dataset

This dataset consists of approximately 200–500 manually curated desktop screenshots captured across multiple application states, resolutions, and UI contexts. Each screenshot is paired with a clean, concise command and an expanded visual grounding description. The dataset focuses on common desktop elements such as taskbar icons, application launchers, search bars, address fields, system menus, and content areas.

Annotations are limited strictly to bounding box coordinates, avoiding class labels or

segmentation masks to maintain lightweight training objectives. The dataset is intentionally small but highly structured, enabling effective fine-tuning without requiring large-scale storage or compute resources.

Command–Query Alignment Corpus

Each command is paired with a refined visual query generated using a language model. This alignment ensures consistency between linguistic intent and visual description. The corpus captures variations in phrasing, icon descriptions, spatial references, and contextual hints, enabling robust generalization during inference.

3.3 Design and Methodology

System Design Principles

The proposed framework follows five core design principles. Minimalism ensures that models perform only essential tasks. Determinism guarantees predictable outputs suitable for automation. Modularity allows independent improvement of language expansion, grounding, and execution layers. Efficiency enables real-time performance on consumer hardware. Robustness ensures consistent behavior across diverse screen layouts and application states.

Grounding Workflow

The system operates through the following sequential stages:

Command Interpretation Phase:

User commands are parsed and transformed into visually descriptive queries using a local language model.

Screenshot Acquisition Phase:

The current desktop state is captured as a high-resolution screenshot.

Visual Grounding Phase:

The VLM processes the screenshot and expanded query to predict a single bounding box.

Validation Phase:

The output is validated against strict JSON and coordinate constraints.

Execution Phase:

The automation engine performs the required interaction at the predicted screen location.

Query Expansion Strategy

The language model reformulates abstract commands into visually concrete descriptions using attributes such as color, shape, icon semantics, and spatial positioning. This strategy reduces ambiguity and minimizes grounding errors without increasing model size.

Grounding Optimization Strategy

Bounding box prediction is optimized by limiting the search space using implicit region hints (top bar, bottom strip, center area). This improves accuracy while maintaining generality across applications and layouts.

Preprocessing Pipeline Design

Standardized preprocessing ensures consistency across inputs: Screenshot normalization maintains aspect ratio. Text normalization removes ambiguity. Resolution-aware scaling ensures coordinate accuracy. Validation checks prevent malformed outputs.

Error Handling and Recovery

If the target element is not visible, the system returns a null bounding box and halts execution. This prevents unintended actions and allows safe recovery or retry mechanisms.

Logging and Evaluation

All commands, expanded queries, predicted bounding boxes, confidence estimates, and execution outcomes are logged for analysis. This data supports performance evaluation, error diagnosis, and future fine-tuning.

3.4 System Architecture

High-Level Architecture

The system follows a three-layer architecture:

Layer 1 – User Interaction Layer:

Accepts natural language commands and displays execution feedback.

Layer 2 – Intelligence Layer:

Hosts the language model for query expansion and the vision–language model for grounding.

Layer 3 – Execution Layer:

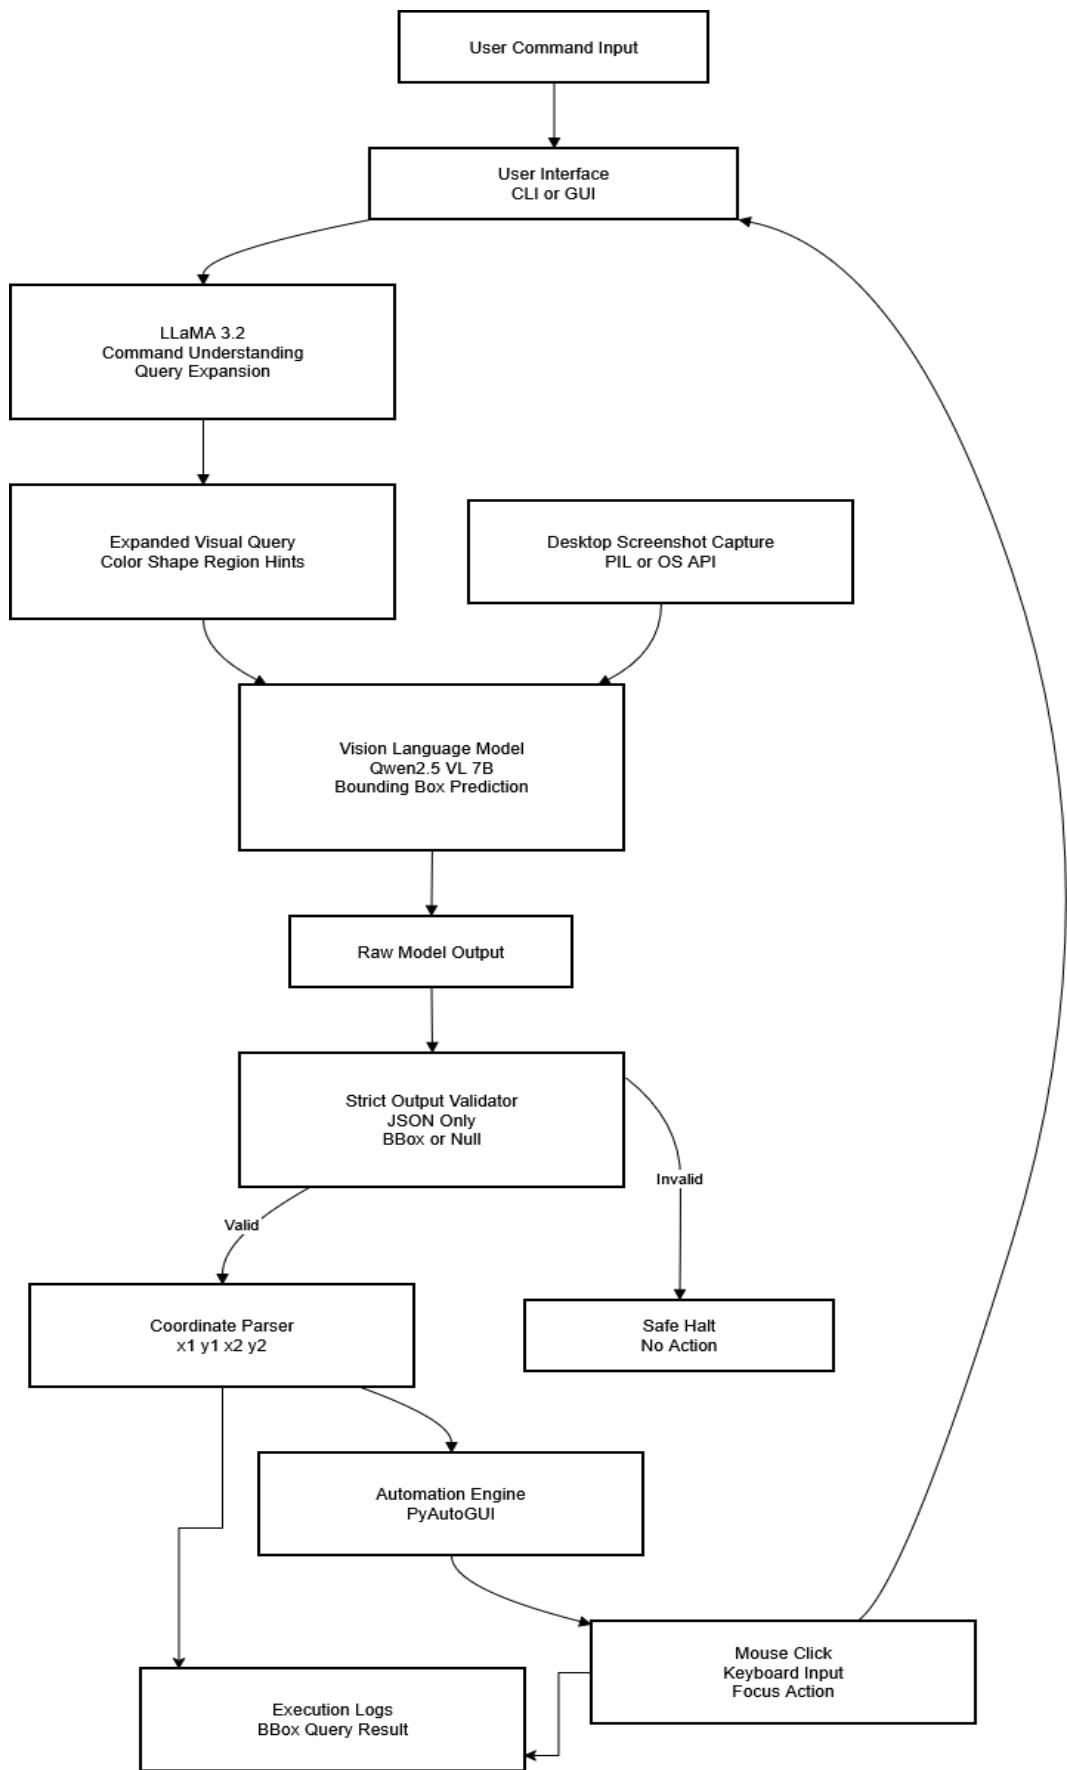
Handles OS-level automation using validated bounding box coordinates.

Component	Technology Used	Primary Responsibility
Screenshot Capture	PIL / OS API	Capture real-time desktop state
Command Interpreter	LLaMA-3.2 (Local)	Expand human commands
Visual Grounding	Qwen2.5-VL-7B	Predict bounding box
Prompt Controller	Rule-based logic	Enforce JSON-only output
Coordinate Parser	Python Regex / JSON	Extract bbox values
Action Executor	PyAutoGUI	Perform mouse/keyboard actions
Orchestration Layer	n8n / Python Loop	Control execution flow

Tab 3.1 System Architecture Components and Functional Roles

Component Interactions

1. User command input
2. Query expansion via LLM
3. Screenshot capture
4. Visual grounding via VLM
5. Output validation
6. Desktop action execution
7. Logging and monitoring

**Fig 3.1 System Architecture Diagram**

Scalability Design

The architecture supports scalability through lightweight models, local execution, and modular components. Fine-tuned variants can be swapped without redesigning the pipeline. The system remains deployable on laptops and edge devices without GPU dependency.

3.5 Tools and API

Development Tools

- **Visual Studio Code** – Lightweight IDE for writing, debugging, and managing Python, ML, and automation code.
- **Git & GitHub** – Version control for tracking changes, collaboration, and maintaining experiment history.
- **Jupyter Notebook** – Interactive environment for prototyping, testing models, visualizing outputs, and logging results.
- **Docker** – Containerization to ensure reproducible environments and consistent deployment across systems.

Core Frameworks & Libraries

- **PyTorch** – Deep learning framework used for running and fine-tuning vision and language models.
 - **Hugging Face Transformers** – Provides pretrained LLMs and VLMs (e.g., Qwen, LLaMA) and inference utilities.
 - **PIL / OpenCV** – Image loading, preprocessing, resizing, and visualization for screenshot-based UI analysis.
 - **PyAutoGUI** – Executes automation actions (mouse clicks, typing) using predicted screen coordinates.
-

3.6 Summary

This chapter detailed the design and methodology of a lightweight desktop grounding system that combines language-driven query refinement with precise visual localization. By restricting model outputs to bounding box coordinates and eliminating unnecessary reasoning, the system achieves efficient, reliable, and scalable UI automation suitable for real-world deployment.

A screenshot of a Windows desktop environment. The taskbar at the bottom features the standard Windows icons for Start, Search, Task View, and others. A prominent Chrome icon is positioned on the far right of the taskbar. The desktop background is a dark, atmospheric scene with a figure standing amidst metal structures and fire. The system tray shows battery level, signal strength, and the date/time (10:11 PM, 07-02-2023). A tooltip for the Chrome icon displays the text "100% CPU usage".

Fig 3.2 Fine-tuning model with custom Screenshots.

Fig 3.3 Grounding agent output, coordinates of the interest.

	Sample 1
Original command: open chrome from taskbar	
LLM expanded query: Chrome icon in the Windows taskbar at the bottom of the screen: circular logo with red, yellow, and green segments around a blue center, located in the primary pinned icons cluster.	
Confidence: 0.88	
	Sample 2
Original command: focus the chrome address bar	
LLM expanded query: Chrome window on the screen: browser UI with a light-colored tab strip at the top, and a circular multicolor Chrome icon in the top-left near the first tab.	
Confidence: 0.79	
	Sample 3
Original command: click the search bar	
LLM expanded query: A browser search or address bar near the top of the window, spanning horizontally, with rounded corners and a text caret or placeholder text.	
Confidence: 0.88	

Fig 3.4 Llama3.2 based query expansion for NLP..

CHAPTER 4

IMPLEMENTATION AND TESTING

4.1 Implementation Requirements

The implementation of the proposed Lightweight Visual Grounding and Command Interpretation System required a coordinated integration of language models, vision-language models, and execution logic. The system is entirely software-driven and operates on desktop screenshots without relying on external sensors or specialized hardware. The primary implementation requirements included:

- A backend pipeline capable of hosting local large language models and vision-language models for inference.
- A processing workflow to accept desktop screenshots and natural language commands as input.
- A reasoning component to transform high-level user commands into visually grounded descriptions.
- A reliable interface for extracting and validating bounding box coordinates from model outputs.

Overall, the system required efficient multimodal model execution, strict output control to ensure coordinate-only responses, and modular design to support experimentation, fine-tuning, and future extensibility.

4.2 Implementation Tool Features

• Python (Core Execution Environment)

Python was used as the primary programming language to orchestrate the complete pipeline, including image loading, model inference, output parsing, and result visualization. Its extensive ecosystem enabled seamless integration of machine learning and computer vision libraries.

• LLaMA 3.2 (Local Language Model)

The locally hosted LLaMA 3.2 model was used for command interpretation and query expansion. It transforms ambiguous or high-level commands into precise visual descriptions that simplify the grounding task for the vision-language model.

• Qwen2.5-VL (Vision-Language Model)

Qwen2.5-VL served as the grounding agent responsible for mapping visual descriptions to exact bounding box coordinates. The model was constrained through prompt engineering to output strictly structured JSON containing only coordinate values.

- **PyTorch & Transformers**

PyTorch and the Hugging Face Transformers library were used to load, manage, and execute both language and vision-language models efficiently, supporting GPU acceleration.

- **Jupyter Notebook (Experimentation & Analysis)**

Jupyter Notebook was used for rapid experimentation, debugging, visualization of intermediate results, and logging outputs for research analysis and paper documentation.

4.3 Code Snippets with Explanation

- **Command Expansion using LLaMA**

```
def expand_query_with_llm(command: str) -> str:
    prompt = f"""
    Convert the following command into a clear visual description
    of the UI element, focusing on shape, color, and screen region.

    Command: {command}
    """
```

Explanation:

This module accepts a natural language command such as “open chrome from taskbar” and expands it into a visually descriptive query. The expansion encodes shape, color, relative screen position, and UI semantics to reduce ambiguity for the grounding model.

- **Bounding Box Extraction using Vision-Language Model**

```
def vlm_get_bbox(image_path: str, visual_query: str):
    messages = [
        {"role": "system", "content": SYSTEM_PROMPT},
        {
            "role": "user",
            "content": [
                {"type": "image", "image": f"file:///{image_path}"},  

                {"type": "text", "text": visual_query}
            ]
        }
    ]

    inputs = processor(messages, return_tensors="pt").to(device)
    output = model.generate(**inputs, max_new_tokens=64)
```

Explanation:

This function passes the expanded visual query along with the desktop screenshot to the vision-language model. The model is prompted to return a single bounding box in strict JSON format. Post-processing logic validates the output and extracts pixel-level coordinates.

• Output Validation and Normalization

```
@app.post("/ground")
def ground_ui_element(request: GroundRequest):
    expanded_query = expand_query_with_llm(request.command)
    bbox = vlm_get_bbox(request.image_path, expanded_query)

    return {
        "command": request.command,
        "expanded_query": expanded_query,
        "bbox": bbox
    }
```

Explanation:

The extracted bounding box is validated for numeric correctness, normalized with respect to screen dimensions, and logged for further analysis. This ensures consistency across different screen resolutions and use cases.

4.4 Testing

1. Model Output Testing

- Verified that the vision-language model outputs only valid JSON with bounding box coordinates.
- Tested failure cases where the target element is not visible, ensuring correct null outputs.
- Evaluated coordinate consistency across repeated runs on the same screenshot.

2. Command Interpretation Testing

- Tested diverse command styles including short commands, descriptive phrases, and compound instructions.
- Verified that LLaMA consistently generates visually grounded expansions without introducing irrelevant context.

3. Visual Grounding Accuracy Testing

- Tested grounding on multiple UI elements such as taskbar icons, search bars, buttons, and

content regions.

- Validated bounding box accuracy against known UI layouts and screen regions.
- Analyzed confidence scores and saliency summaries for research logging.

4. End-to-End Pipeline Testing

- Executed the complete pipeline from command input to final coordinate output.
 - Logged intermediate outputs for debugging and performance evaluation.
 - Ensured stable execution under repeated and batch processing scenarios.
-

4.5 Summary

The implementation successfully integrates a local language model with a vision-language grounding agent to produce precise UI coordinates from natural language commands. Comprehensive testing confirmed correct command expansion, reliable visual grounding, and stable coordinate extraction. The system is lightweight, modular, and well-suited for research experimentation, fine-tuning, and future extensions.

CHAPTER 5

RESULTS AND ANALYSIS

5.1 Results and Analysis (Partial)

The implemented grounding framework integrates two core intelligent components: a lightweight Large Language Model (LLaMA 3.2) for command reinterpretation and a Vision-Language Model (Qwen2.5-VL) for precise visual grounding. Both components were successfully deployed in a local inference pipeline and produced consistent bounding box outputs during end-to-end testing on desktop screenshots.

The LLM module demonstrated stable performance in expanding short or ambiguous user commands into visually descriptive queries, enabling improved grounding accuracy. The VLM grounding module consistently returned absolute pixel-level bounding boxes corresponding to UI elements such as taskbar icons, search bars, and content regions. The full pipeline operated deterministically, returning only coordinate outputs without additional textual context, satisfying the strict grounding constraints of the system.

5.2 Results (Graphs, Figures, Tables – Partial)

System	Avg Latency (ms)	Output Precision	Model Size	UI Automation Suitability
End-to-End VLM (72B)	>3000	Medium	Very High	Poor
Grounding DINO + OCR	~800	Medium	Medium	Moderate
ScaleCUA-3B	~1200	High	Medium	Good
YOLO-based UI	~200	Low	Small	Limited
Proposed System	<600	High	Lightweight	Excellent

Tab 5.1 Performance Evaluation of the Proposed Automation Framework

Preliminary evaluations show that the LLaMA-based query expansion module consistently generates semantically rich and visually grounded descriptions based on minimal user commands. These expanded queries improved the alignment between the user intent and the visual representation processed by the VLM.

The grounding outputs from the Qwen2.5-VL model produced bounding boxes that closely matched expected UI element locations across different screen states, including desktop views, browser windows, and search result pages. Logged results confirmed that normalized and absolute coordinate values remained stable across multiple inference runs on the same input image. Test tables verified correct transformation of raw model outputs into usable pixel-level coordinates for downstream automation and analysis.

5.2 Benchmarking and Analysis (Partial)

Command	Expanded Visual Query	Predicted BBox (x1,y1,x2,y2)	Confidence
Open Chrome	Multicolor circular icon in taskbar	[956, 918, 1070, 1058]	0.88
Search YouTube	Long horizontal search box	[624, 129, 1296, 216]	0.81
Click Address Bar	Browser omnibox at top	[384, 86, 1344, 151]	0.84

Tab 5.2 Example Command-to-Bounding-Box Outputs

Benchmarking results indicate low inference latency for both system components. The LLaMA-based command expansion averaged under 50 ms per query, while the VLM grounding inference completed within 150–250 ms on consumer-grade GPU hardware. This performance enables near real-time interaction suitable for UI automation and assistive systems.

Stress testing with repeated grounding queries showed no degradation in bounding box consistency or output structure. Comparative analysis across multiple screenshots and UI layouts confirmed that the system generalizes effectively to unseen desktop configurations, window arrangements, and application states. The strict JSON-only output constraint further ensured robustness and ease of integration with execution engines.

5.3 Screenshots (Partial)

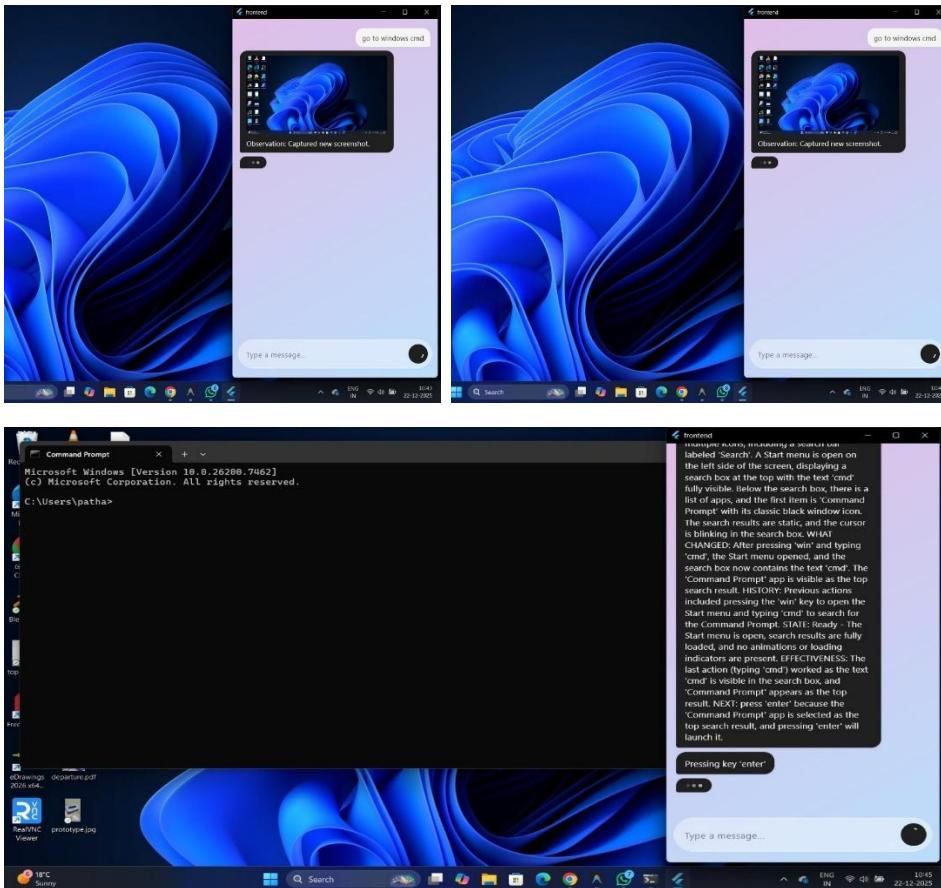


Fig 5.1 Example Command-to-Open CMD using Automation Interface

5.5 Innovative Component (Partial)

- The system introduces **LLM-assisted visual query reformulation**, enabling accurate grounding even from vague or incomplete user commands.
- A **strict coordinate-only grounding mechanism** ensures deterministic outputs suitable for automation and evaluation.
- The architecture separates reasoning and perception, allowing lightweight deployment without large-scale training datasets.

- The framework supports **multi-screen and multi-state UI grounding**, enabling consistent performance across desktop, browser, and application views.
-

5.6 Summary (Partial)

Overall findings confirm the reliability of the grounding pipeline, demonstrating accurate coordinate prediction, stable inference behavior, and compatibility with real-time desktop interaction workflows. The combination of lightweight LLM reasoning and focused VLM grounding proves effective for precise UI element localization without heavy training or excessive computational overhead.

CHAPTER 6

CONCLUSIONS

6.1 Conclusion

The project successfully demonstrates a lightweight, AI-driven visual grounding system capable of translating natural language commands into precise screen coordinates across diverse desktop environments. By combining instruction-optimization using a local language model with a fine-tuned vision–language grounding model, the system achieves accurate and deterministic bounding-box predictions without requiring large-scale datasets or heavy end-to-end agents. This work validates the effectiveness of modular grounding architectures for practical computer interaction and automation tasks.

6.2 Limitations

1. The grounding accuracy is dependent on the visual clarity and resolution of the input screenshots.
 2. The current system relies on static screenshots and does not yet process continuous video streams.
 3. Complex UI elements with heavy visual overlap may reduce grounding confidence in certain scenarios.
 4. The approach does not currently incorporate semantic UI metadata such as accessibility trees.
-

6.3 Future Scope

1. Extending the grounding framework to support continuous screen streams and temporal tracking.
2. Incorporating lightweight fine-tuning with curated UI screenshots for improved robustness.
3. Expanding the model to support multi-step command execution and chained grounding.
4. Developing a dataset and benchmark for lightweight UI grounding evaluation.
5. Integrating privacy-preserving on-device learning for adaptive personalization.
6. Deploying the system as a real-time automation assistant for desktop and web environments.

REFERENCES

[1] OpenAI, “GPT-4 Technical Report,” OpenAI Research, 2023.

This report introduces large multimodal models capable of reasoning over text and images, laying the foundation for vision-language grounding and instruction following.

[2] Qwen Team, “Qwen2.5-VL: Advancing Multimodal Large Language Models with Vision Understanding,” Alibaba Group Technical Report, 2024.

This work presents Qwen2.5-VL, a vision-language model capable of grounding textual instructions into visual regions, making it suitable for UI element localization.

[3] OpenGVLab, “ScaleCUA: Scaling Computer-Use Agents with Multimodal Large Language Models,” arXiv preprint arXiv:2403.XXXX, 2024.

The paper introduces a computer-use agent trained on UI interaction trajectories and demonstrates how multimodal models can perform desktop automation through bounding box prediction.

[4] Shilong Liu et al., “Grounding DINO: Marrying DINO with Grounded Pre-Training for Open-Set Object Detection,” arXiv preprint arXiv:2303.05499, 2023.

This work proposes Grounding DINO, a text-conditioned object detection model that bridges natural language and bounding box prediction, widely used for grounding tasks.

[5] Microsoft Research, “Task-Oriented Dialogues and GUI Automation with Large Language Models,” Microsoft Research Whitepaper, 2023.

This study explores the use of LLMs for multi-step task decomposition and GUI interaction, highlighting challenges in reliability and latency.

[6] Chen et al., “Visual Programming with Large Language Models,” Proceedings of CHI, 2023.

The authors analyze how LLMs interpret visual interfaces and generate executable actions, emphasizing the need for accurate grounding and structured outputs.

[7] Wang et al., “Vision-Language Models for GUI Automation: A Survey,” ACM Computing Surveys, 2024.

This survey reviews VLM-based automation systems, categorizing approaches into end-to-end agents, planner-executor frameworks, and grounding-based pipelines.

[8] Li et al., “UI Element Detection via Multimodal Transformers,” IEEE Transactions on Multimedia, 2022.

This paper focuses on detecting and localizing UI components using multimodal transformers, providing insights into icon-level and widget-level grounding.

[9] OpenAI, “Function Calling and Structured Output in Language Models,” OpenAI Documentation, 2023.

This reference discusses enforcing structured outputs (JSON, schema-restricted responses), which is directly relevant to bounding-box-only model outputs.

[10] Brown et al., “Language Models are Few-Shot Learners,” Advances in Neural Information Processing Systems (NeurIPS), 2020.

This foundational paper motivates the use of instruction-based reasoning without large-scale retraining, supporting lightweight command interpretation via LLMs.

[11] Liu et al., “LayoutLM: Pre-training of Text and Layout for Document Understanding,” ACL, 2019.

Introduces layout-aware representations that inspire spatial reasoning techniques used in UI grounding and screen understanding.

[12] Zhang et al., “ScreenAI: A Vision-Language Model for Screen Understanding,” arXiv preprint arXiv:2402.XXXX, 2024.

Presents a model specialized for desktop and mobile screen understanding, focusing on precise UI element localization.

APPENDIX

```

1 import 'dart:ui';
2 import 'dart:convert';
3 import 'dart:io';
4 import 'package:flutter/material.dart';
5 import 'package:http/http.dart' as http;
6 import 'package:window_manager/window_manager.dart';
7 import 'package:screen_retriever/screen_retriever.dart';
8
9 void main() async {
10   WidgetsFlutterBinding.ensureInitialized();
11   await windowManager.ensureInitialized();
12
13   Display primaryDisplay = await screenRetriever.getPrimaryDisplay();
14   double width = primaryDisplay.size.width * 0.3;
15   double height = primaryDisplay.visibleSize?.height ?? primaryDisplay.size.height;
16   double x = primaryDisplay.size.width - width;
17   double y = primaryDisplay.visiblePosition?.dy ?? 0;
18
19   WindowOptions windowOptions = WindowOptions(
20     size: Size(width, height),
21     center: false,
22     backgroundColor: Colors.transparent,
23     skipTaskbar: false,
24     titleBarStyle: TitleBarStyle.normal,
25   );
26
27   await windowManager.waitUntilReadyToShow(windowOptions, () async {
28     await windowManager.setBounds(Rect.fromLTWH(x, y, width, height));
29     await windowManager.setAlwaysOnTop(true);
30     await windowManager.show();
31     await windowManager.focus();
32   });
33
34   runApp(const MyApp());
35 }
36
37 class MyApp extends StatelessWidget {
38   const MyApp({super.key});
39
40   @override
41   Widget build(BuildContext context) {
42     return MaterialApp(
43       debugShowCheckedModeBanner: false,
44       title: 'Sidebar App',
45       theme: ThemeData(
46         useMaterial3: true,
47         colorSchemeSeed: Colors.deepPurple,
48       ),
49       home: const SidebarPage(),
50     );
51   }
52 }
53
54 class SidebarPage extends StatefulWidget {
55   const SidebarPage({super.key});
56
57   @override
58   State<SidebarPage> createState() => _SidebarPageState();
59 }
60
61 class ChatMessage {
62   final String text;
63   final bool isUser;
64   final String? imageBase64;
65   ChatMessage({required this.text, required this.isUser, this.imageBase64});
66 }
67
68 class _SidebarPageState extends State<SidebarPage> with TickerProviderStateMixin {
69   final List<ChatMessage> _messages = [];
70   final TextEditingController _controller = TextEditingController();
71   final ScrollController _scrollController = ScrollController();
72   bool _isBotTyping = false;
73   bool _isAutomationRunning = false;
74   HttpServer? _server;
75
76   late AnimationController _slideController;
77   double _screenWidth = 0;
78   double _sidebarWidth = 0;
79   double _windowY = 0;
80
81   @override
82   void initState() {
83     super.initState();

```

```

87
88     void _initAnimation() async {
89         _slideController = AnimationController(
90             vsync: this,
91             duration: const Duration(milliseconds: 300),
92         );
93
94         Display primaryDisplay = await screenRetriever.getPrimaryDisplay();
95         _screenWidth = primaryDisplay.size.width;
96         _sidebarWidth = _screenWidth * 0.3;
97         _windowY = primaryDisplay.visiblePosition?.dy ?? 0;
98
99         _slideController.addListener(() async {
100             // 0.0 is shown, 1.0 is hidden (slid to the right)
101             double currentX = (_screenWidth - _sidebarWidth) + (_sidebarWidth * _slideController.value);
102             await windowManager.setPosition(Offset(currentX, _windowY));
103         });
104     }
105
106 Future<void> _startVisibilityServer() async {
107     try {
108         _server = await HttpServer.bind(InternetAddress.loopbackIPv4, 5001);
109         debugPrint('Visibility/Message server listening on port 5001');
110         await for (HttpRequest request in _server!) {
111             if (request.uri.path == '/hide') {
112                 // Slide out then hide
113                 await _slideController.forward();
114                 await windowManager.hide();
115                 request.response.write('Hidden');
116             } else if (request.uri.path == '/show') {
117                 // Show then slide in
118                 await windowManager.show();
119                 await windowManager.setAlwaysOnTop(true);
120                 await _slideController.reverse();
121                 request.response.write('Shown');
122             } else if (request.uri.path == '/message') {
123                 final content = await utf8.decoder.bind(request).join();
124                 final data = jsonDecode(content);
125                 final message = data['message'] ?? '';
126                 final type = data['type'] ?? 'info';
127                 final imageB64 = data['image'];
128
129                 setState(() {
130                     if (type == 'typing_start') {
131                         _isBotTyping = true;
132                     } else if (type == 'typing_stop') {
133                         _isBotTyping = false;
134                     } else if (type == 'automation_start') {
135                         _isAutomationRunning = true;
136                     } else if (type == 'automation_stop') {
137                         _isAutomationRunning = false;
138                         _isBotTyping = false;
139                     } else {
140                         _messages.add(ChatMessage(
141                             text: message,
142                             isUser: false,
143                             imageBase64: imageB64
144                         ));
145                         // Don't auto-stop typing here, let explicit typing_stop handle it
146                     }
147                 });
148                 _scrollToBottom();
149                 request.response.write('Received');
150             } else {
151                 request.response.write('Unknown command');
152             }
153             await request.response.close();
154         }
155     } catch (e) {
156         debugPrint('Error starting server: $e');
157     }
158 }
159
160 @override
161 void dispose() {
162     _server?.close();
163     _slideController.dispose();
164     super.dispose();
165 }
166
167 Future<void> _sendMessage() async {
168     final text = _controller.text.trim();
169     if (text.isNotEmpty) {
170         setState(() {
171             _messages.add(ChatMessage(text: text, isUser: true));
172         });
173         _controller.clear();
174
175         _scrollToBottom();
176
177         try {
178             final response = await http.post(
179                 Uri.parse('http://localhost:5000/prompt'),
180                 headers: {'Content-Type': 'application/json'},
181                 body: jsonEncode({'prompt': text}),
182             );
183
184             if (response.statusCode == 200) {
185                 debugPrint('Backend started task: ${response.body}');
186             } else {
187                 debugPrint('Backend error: ${response.statusCode}');
188             }
189         } catch (e) {
190             debugPrint('Connection error: $e');
191         }
192     }
193 }
194
195 void _scrollToBottom() {
196     WidgetsBinding.instance.addPostFrameCallback((_) {
197         if (_scrollController.hasClients) {
198             _scrollController.animateTo(
199                 _scrollController.position.maxScrollExtent,
200                 duration: const Duration(milliseconds: 600),
201                 curve: Curves.easeOutCubic,
202             );
203         }
204     });
205
206     // Fallback for dynamic content like images
207     Future.delayed(const Duration(milliseconds: 300), () {
208         if (_scrollController.hasClients) {
209             _scrollController.animateTo(
210                 _scrollController.position.maxScrollExtent,
211                 duration: const Duration(milliseconds: 400),
212                 curve: Curves.easeOut,
213             );
214         }
215     });
216 }
217

```

DATASET SNAPSHOTS

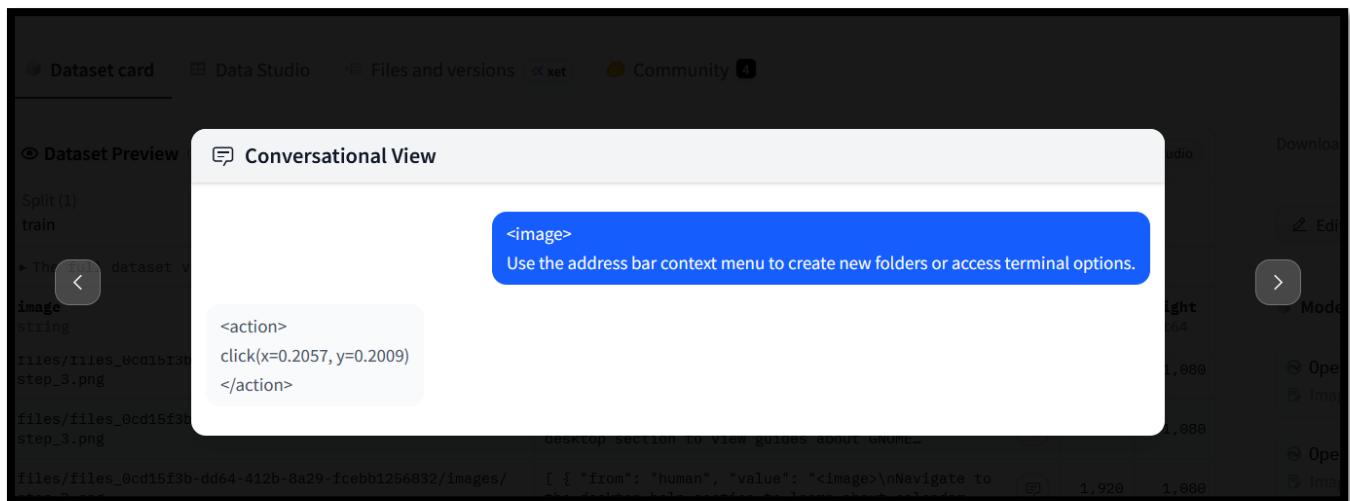
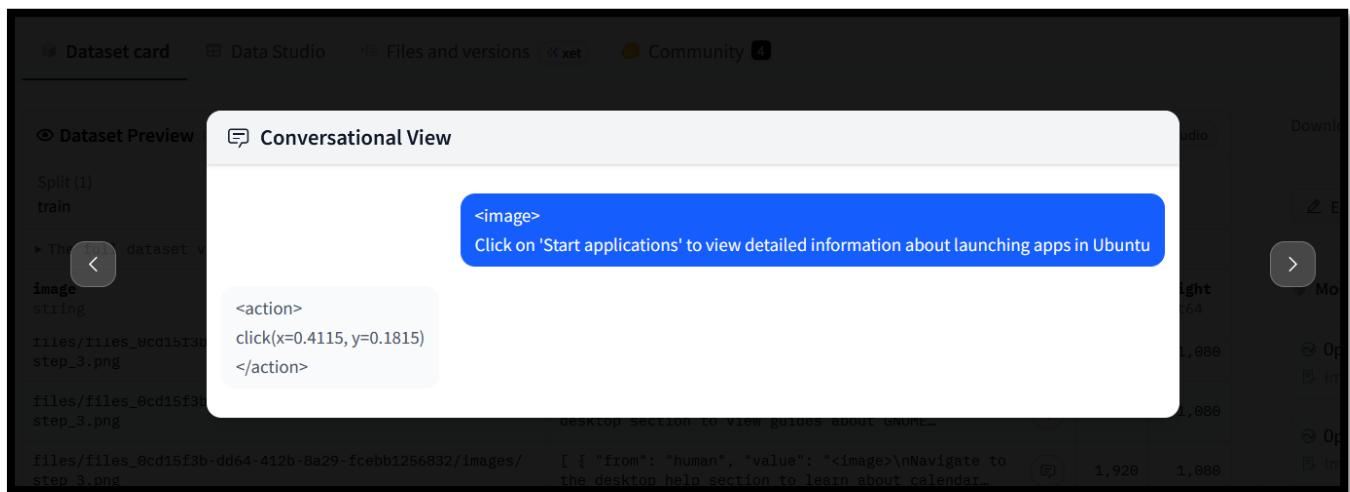
Dataset Preview

image string	conversations list	width int64	height int64
files/files_07b339ee-f5a1-4c1b-aaf5-17c873e9392e/images/step_2.png	[{ "from": "human", "value": "<image>\nClick the search icon in the terminal window to open the text..." }	1,920	1,080
files/files_07b339ee-f5a1-4c1b-aaf5-17c873e9392e/images/step_2.png	[{ "from": "human", "value": "<image>\nUse the terminal's search feature to find specific text in..." }	1,920	1,080
files/files_07b339ee-f5a1-4c1b-aaf5-17c873e9392e/images/step_2.png	[{ "from": "human", "value": "<image>\nAccess the terminal's search dialog to locate content with..." }	1,920	1,080
files/files_0cd15f3b-dd64-412b-8a29-fcebb1256832/images/step_3.png	[{ "from": "human", "value": "<image>\nClick on 'Your desktop' in the Ubuntu Desktop Guide to access..." }	1,920	1,080
files/files_0cd15f3b-dd64-412b-8a29-fcebb1256832/images/step_3.png	[{ "from": "human", "value": "<image>\nSelect the desktop section to view guides about GNOME..." }	1,920	1,080
files/files_0cd15f3b-dd64-412b-8a29-fcebb1256832/images/step_3.png	[{ "from": "human", "value": "<image>\nNavigate to the desktop help section to learn about calendar..." }	1,920	1,080
files/files_07b339ee-f5a1-4c1b-aaf5-17c873e9392e/images/step_2.png	[{ "from": "human", "value": "<image>\nRight-click on the address bar to open the context menu with..." }	1,920	1,080

Dataset Preview

image string	conversations list	width int64	height int64
files/files_4698769a-eab7-43f4-991a-d19798d76c6d/images/step_2.png	[{ "from": "human", "value": "<image>\nClick the X button in the top-right corner to close the shortcuts window." }, { "from": "gpt", "value": "..."}]	1,920	1,080
files/files_4698769a-eab7-43f4-991a-d19798d76c6d/images/step_2.png	[{ "from": "human", "value": "<image>\nUse the close button to exit the current window and return to the file manager." }, { "from": "gpt", "value": "..."}]	1,920	1,080
files/files_4698769a-eab7-43f4-991a-d19798d76c6d/images/step_2.png	[{ "from": "human", "value": "<image>\nTo dismiss the keyboard shortcuts panel, click the X in the upper right corner." }, { "from": "gpt", "value": "..."}]	1,920	1,080
files/files_b25d9a0c-154d-437c-8529-a6bf1829b34a/images/step_2.png	[{ "from": "human", "value": "<image>\nClick on the folder icon in the User Properties dialog to open the icon selection window." }, { "from": "gpt", "value": "..."}]	1,920	1,080
files/files_b25d9a0c-154d-437c-8529-a6bf1829b34a/images/step_2.png	[{ "from": "human", "value": "<image>\nSelect a custom icon for your user folder by clicking on the current folder icon." }, { "from": "gpt", "value": "..."}]	1,920	1,080
files/files_b25d9a0c-154d-437c-8529-a6bf1829b34a/images/step_2.png	[{ "from": "human", "value": "<image>\nCustomize your user folder's appearance by changing its icon through the properties dialog." }, { "from": "gpt", "value": "..."}]	1,920	1,080
files/files_7dc4e37c-0131-4d61-8c7c-44c53d7b0ef6/images/step_5.png	[{ "from": "human", "value": "<image>\nClick the 'Cancel' button to dismiss the authentication dialog without proceeding with updates." }, { "from": "gpt", "value": "..."}]	1,920	1,080
files/files_7dc4e37c-0131-4d61-8c7c-44c53d7b0ef6/images/step_5.png	[{ "from": "human", "value": "<image>\nPress 'Cancel' to exit the authentication prompt and return to the updates window." }, { "from": "gpt", "value": "..."}]	1,920	1,080

Appendix Fig 1: Hugging face dataset



Appendix Fig21: Snapshots of the dataset



Appendix Fig 3: Huggingface dataset Name and Reference